

STAR NLP CLASSIFICATION PROJECT

Fake or Real? Detecting News with NLP

Presented by:

Eya Cherif, Isis Hassan and Suzana Cracco

Feb.2026



Problem Statement & Data

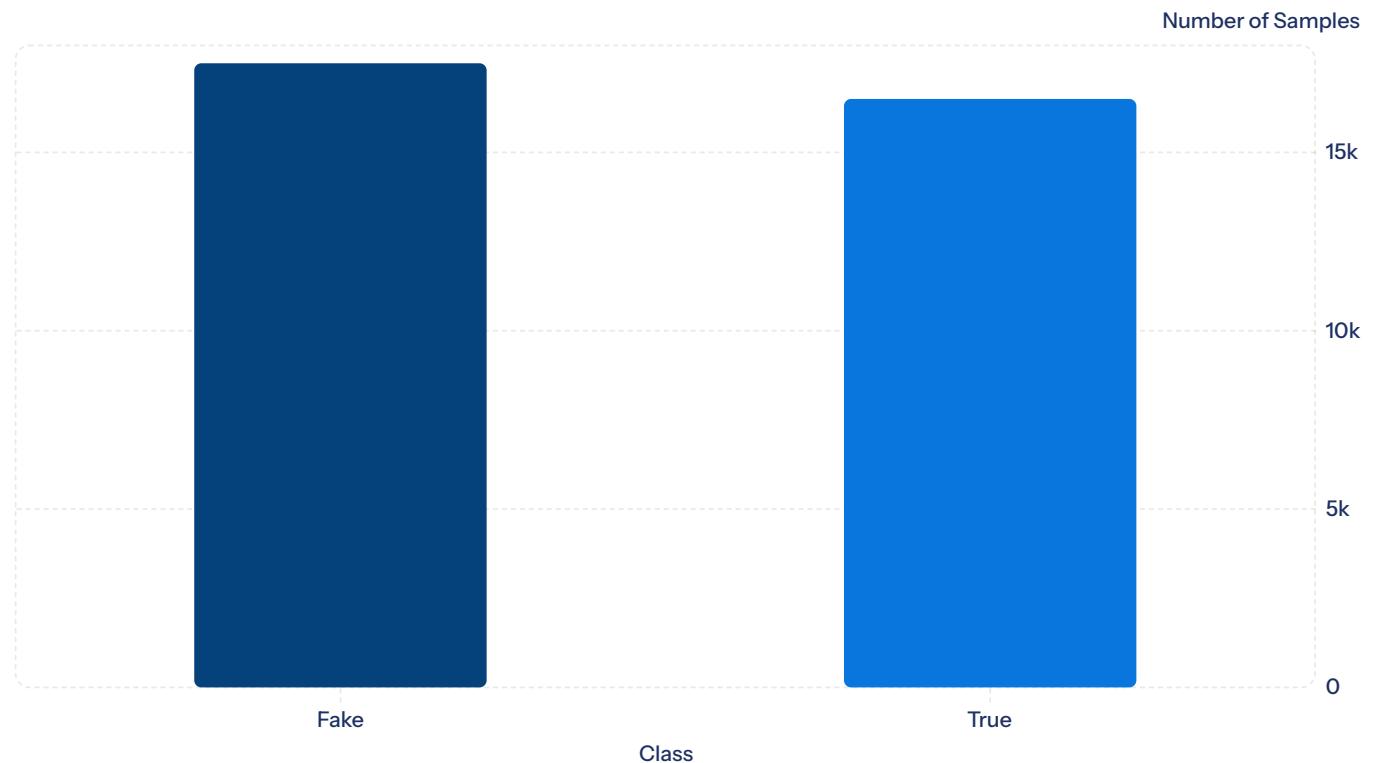
Can machines spot fake news?

Kaggle Fake News Detection dataset:

32,206 labeled articles combined

Custom test set of 100 headlines

50 real (2025 news from CNN, Wikipedia, Britannica) and 50 satirical (The Onion).



The distribution of samples per class shows a good balance in the training data.



⚙️ PREPROCESSING

Data Preprocessing Pipeline

Lemmatization

Optional: reduce words to their root form.

Tokenization

Split cleaned text into individual tokens.



Text Cleaning

Remove punctuation, stop words and lowercase text.

Vectorization

Convert tokens to BoW, TF-IDF, or BERT vectors.

Tokenization Techniques



Bag of Words

Counts word occurrences. Simple but loses word order and context.



TF-IDF

Weighs terms by frequency and rarity. Better at surfacing distinctive words.

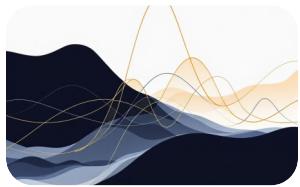


BERT Transformer

Contextual embeddings from a pretrained language model. Captures meaning and nuance.

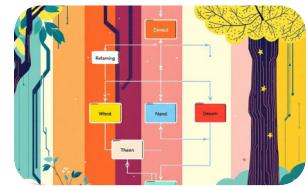


Model Architectures



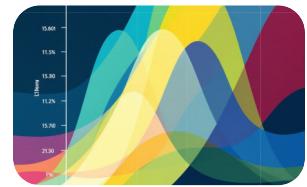
Logistic Regression

Linear baseline. Fast training, interpretable coefficients.



XGBoost

Gradient-boosted trees. Handles sparse text features well.



MultinomialNB

Naïve Bayes for text. Strong with BoW/TF-IDF features.



BERT

Fine-tuned transformer. Learns deep contextual representations.

Validation vs Test Performance

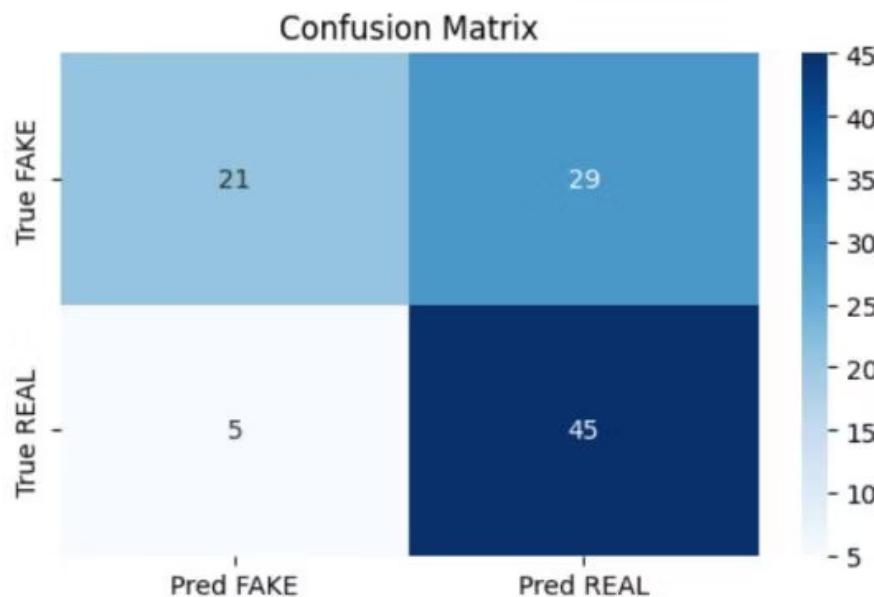
All models achieved **high validation accuracy (90–98%)**, but performance dropped on our custom 2025 test set: **Generalization challenges**

Model	Val Acc (80 Train/20 Test Split)	Test Acc (generated dataset)
TF-IDF + XGB (NoLem)	0.88	0.30
TF-IDF + XGB (Lem)	0.90	0.64
BoW + XGB (NoLem)	0.93	0.42
BoW + XGB (Lem)	0.94	0.39
BERT (bert-base-uncased)	0.98	0.64
BERT + Swear Flag	0.98	0.66
TF-IDF + LogReg	0.93	0.69
BoW + LogReg (tuned)	0.94	0.73
TF-IDF + LogReg	0.93	0.69
Naive Bayes (BoW)	0.96	0.31
Naive Bayes (TF-IDF)	0.98	0.3
DistilBERT (128)	0.78	—
DistilBERT (256)	0.98	—

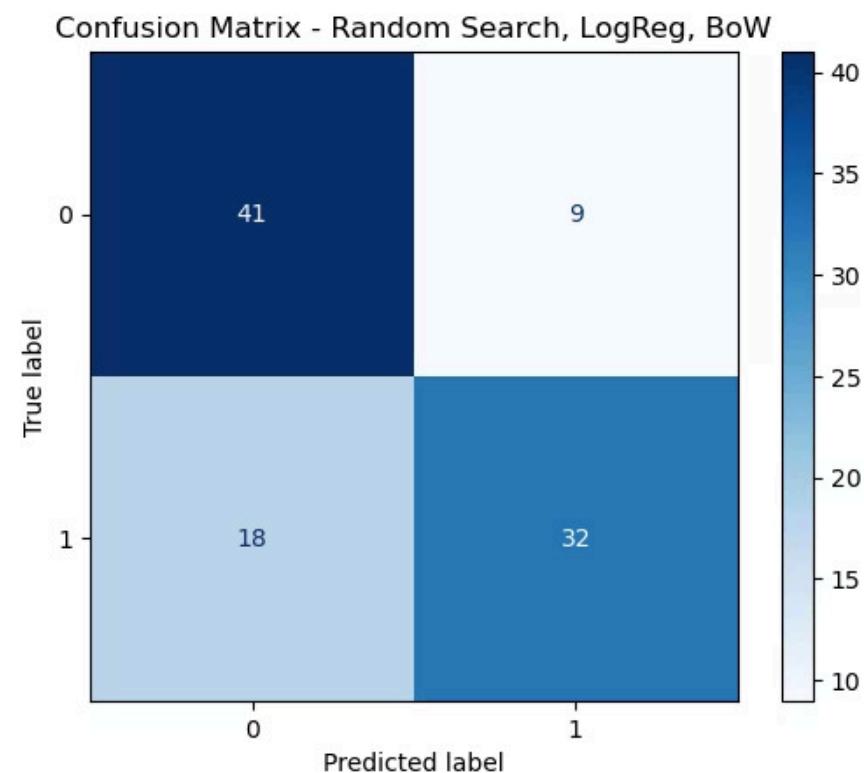
Confusion Matrices – Key Models

Confusion Matrices (Test)

BERT + Swear Flag



LogisticRegression (BoW)



Key Takeaways

1 Validation ≠ Real-World Performance

All models scored 90–98% on validation but dropped significantly on unseen 2025 headlines, revealing the generalization gap.

3 Simplicity Often Wins

Logistic Regression achieved the best test accuracy at 69%, outperforming more complex models.

5 Next Steps

Larger, more diverse training data and domain adaptation techniques could close the generalization gap.

2 Preprocessing Matters (Selectively)

Lemmatization was critical for XGBoost (30% → 64% accuracy), but other models achieved good results with minimal preprocessing.

4 BERT Shows Promise

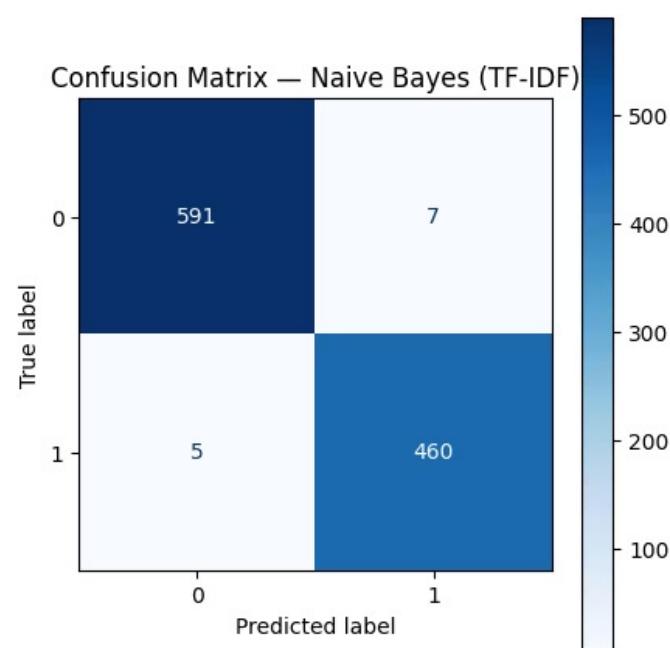
Transformers outperformed traditional methods, and adding a swear word flag improved generalization.

Appendix

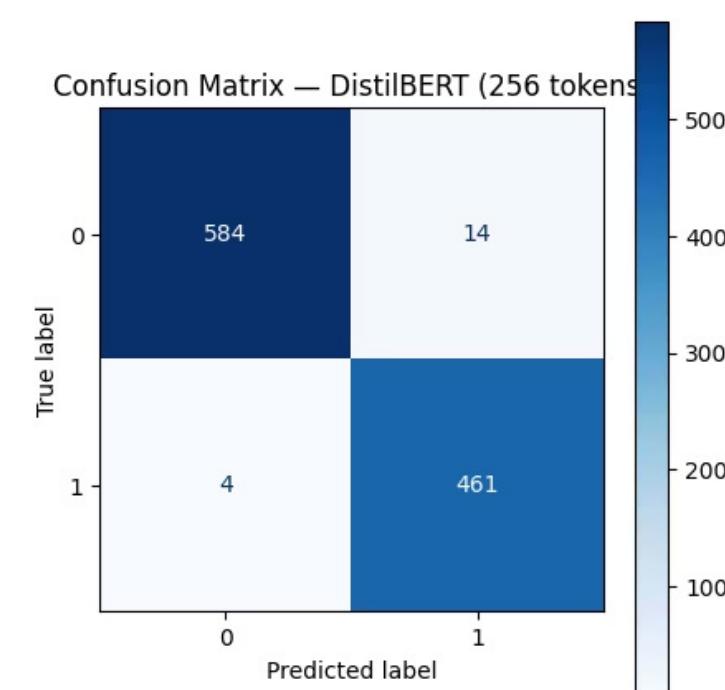
Confusion Matrices – Key Models

Confusion Matrices (Validation)

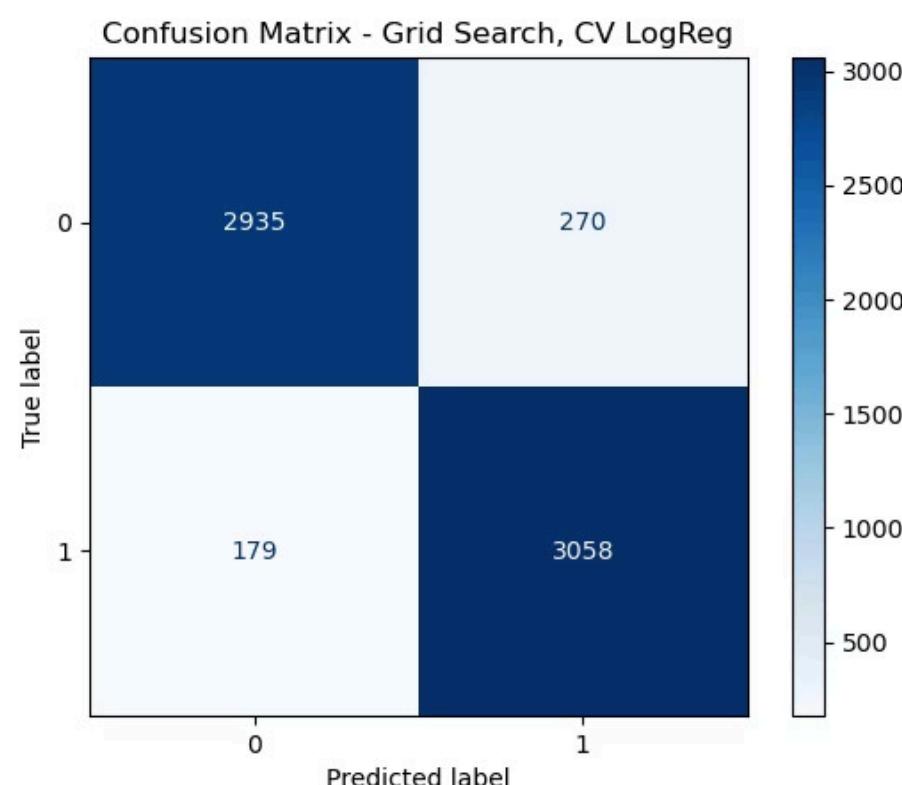
Naive Bayes (TF-IDF)



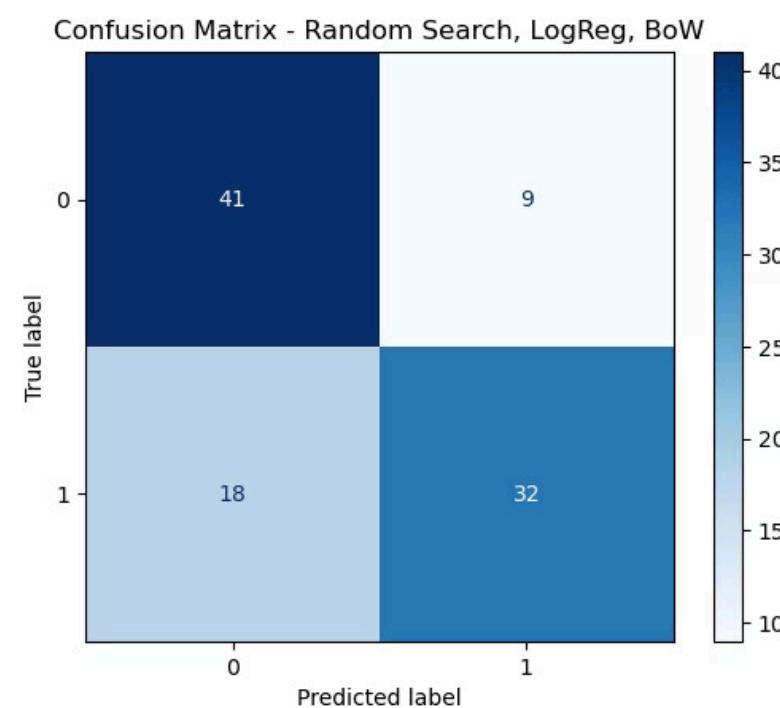
DistilBERT (256 tokens)



LogisticRegression (Bow) - Val

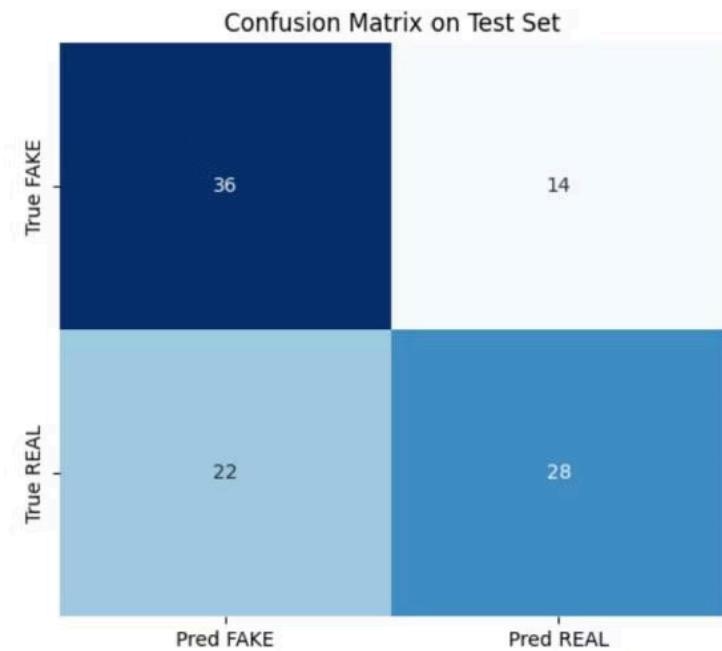


LogisticRegression (Bow) - Test - keep

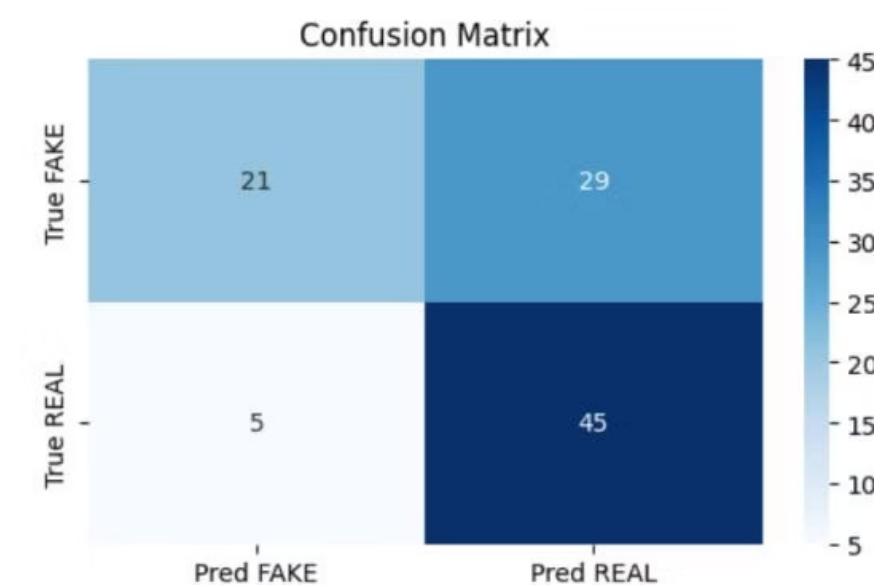


Confusion Matrices (Test)

TF-IDF + XGB (Lem)



BERT + Swear Flag – Test



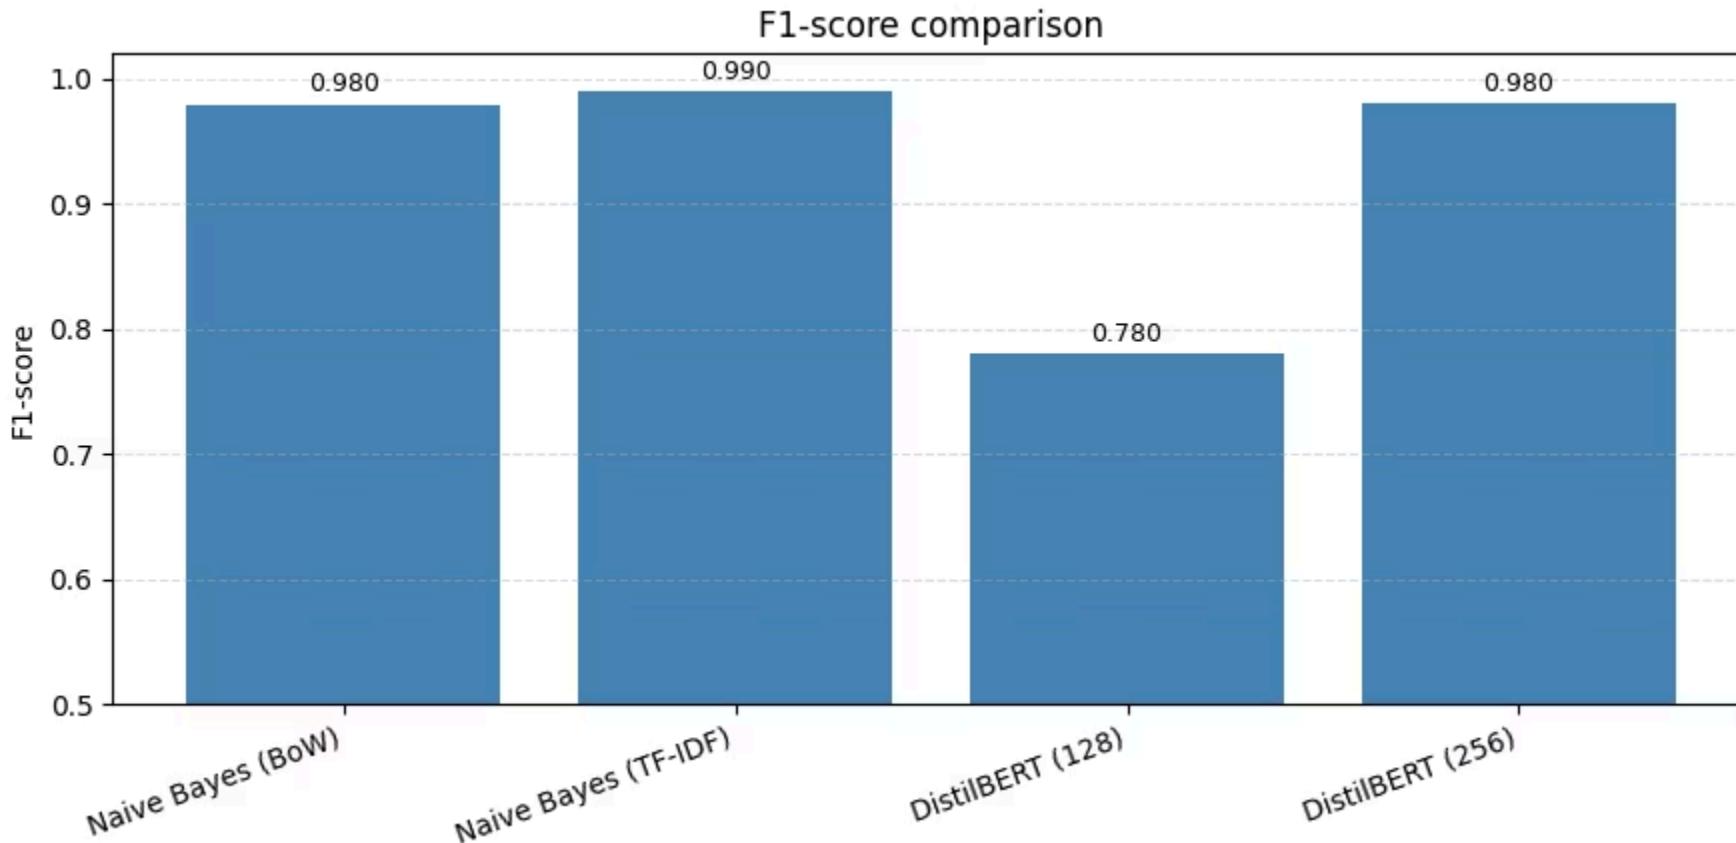
Additional Model Experiments

To explore model performance further, we evaluated additional Naive Bayes and DistilBERT configurations. These results highlight their validation accuracy and class-specific recall before deployment to our custom test set.

Model Performance Metrics (Validation)

Model	Val Acc	F1-Score	Recall (Class 0)	Recall (Class 1)
Naive Bayes (BoW)	0.960	0.960	0.980	0.980
Naive Bayes (TF-IDF)	0.980	0.990	0.980	0.990
DistilBERT (128)	0.780	0.780	0.640	0.980
DistilBERT (256)	0.980	0.980	0.980	0.990

Validation F1-Scores



DistilBERT models, especially with longer sequence lengths (256 tokens), demonstrated superior F1-scores during validation.

Our Custom Test Headlines

Real Headlines (2025)

- Trump inaugurated as 47th President
- LA wildfires burn 50,000+ acres
- Nvidia loses ~\$600B after DeepSeek
- Israel–Hamas ceasefire agreement

Satirical Headlines (The Onion)

- "Man Who Bumped Tesla... Sentenced To Death"
- "Archaeologists Uncover Ancient Race Of Skeleton People"
- "No Way To Prevent This, Says Only Nation..."
- "Zoo Gorilla Looks Bored Reading Wuthering Heights"

50 real headlines from CNN, Wikipedia, Britannica/CBS · 50 fake from The Onion across politics, science, culture.